


Convolutional Autoencoders for Dimensionality Reduction



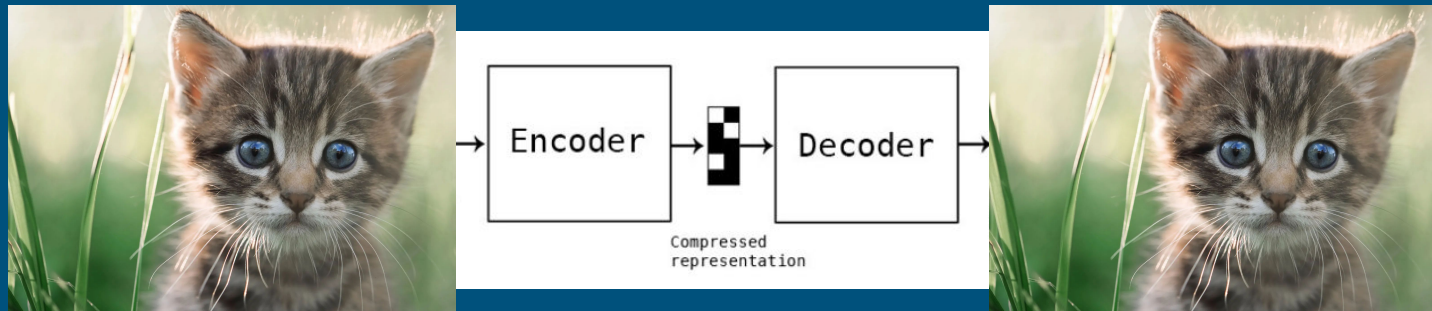
By Luke Harwood | Kevin Paganini



Type of machine learning problem

Main Purpose: Reducing Complexity/Dimensionality of Feature Space

- Unsupervised learning technique
- Used in regression or classification
- Performs Feature Extraction (or Feature Transformation)
- Input and output are the same image



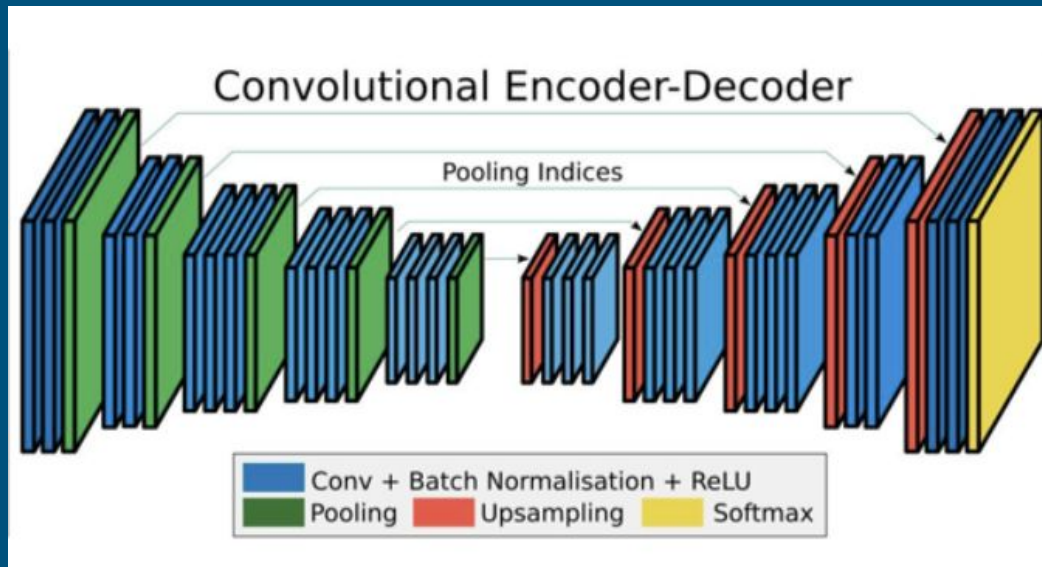
Model Architecture

Model consists of three main parts

- The Encoder
- Latent Layer
- The Decoder

Input Image:
128, 128, 3

Output Image (same):
128, 128, 3



Latent Space:
1024

Train Image Shapes:

Model implementation


Used keras + tensorflow to implement the sequential model

Image preprocessing: CV2 → Resizing images to 128, 128, 3

Encoder: Conv2D, MaxPooling2D, GlobalMaxPooling2D, Dense

Decoder: Conv2D, UpSampling2D, Reshape


Image display: Matplotlib



(150, 150, 3)	13986
(113, 150, 3)	7
(111, 150, 3)	3
(135, 150, 3)	3
(144, 150, 3)	2
(123, 150, 3)	2
(142, 150, 3)	2
(146, 150, 3)	2
(143, 150, 3)	2
(134, 150, 3)	2
(136, 150, 3)	2
(108, 150, 3)	2
(105, 150, 3)	1
(97, 150, 3)	1
(131, 150, 3)	1
(147, 150, 3)	1
(81, 150, 3)	1
(145, 150, 3)	1
(141, 150, 3)	1
(100, 150, 3)	1
(103, 150, 3)	1
(76, 150, 3)	1
(120, 150, 3)	1
(102, 150, 3)	1
(119, 150, 3)	1
(133, 150, 3)	1
(115, 150, 3)	1
(124, 150, 3)	1
(110, 150, 3)	1
(149, 150, 3)	1
(140, 150, 3)	1
dtype: int64	

Encoder Architecture

- Goes from high-dimensional space to low-dimensional space
- 5 Convolutional layers with increasing filter number
- 4 MaxPooling layers
- 1 GlobalMaxPooling
- Latent Layer: 1024 Dense layer



Encoder


Latent Layer

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 32)	896
max_pooling2d (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_1 (Conv2D)	(None, 64, 64, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_2 (Conv2D)	(None, 32, 32, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_3 (Conv2D)	(None, 16, 16, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_4 (Conv2D)	(None, 8, 8, 512)	1180160
global_max_pooling2d (GlobalMaxPooling2D)	(None, 512)	0
dense_6 (Dense)	(None, 1024)	525312
Total params: 2,093,888		
Trainable params: 2,093,888		
Non-trainable params: 0		

Decoder Architecture

- Goes from low-dimensional space to high-dimensional space
- Similar but in reverse
- 1 Reshape layer
- 6 Convolutional layers
- 3 UpSampling layers

Upsampling is the opposite of Pooling



Decoder

Layer (type)	Output Shape	Param #
reshape_1 (Reshape)	(None, 16, 16, 4)	0
conv2d_5 (Conv2D)	(None, 16, 16, 512)	18944
up_sampling2d (UpSampling2D)	(None, 32, 32, 512)	0
conv2d_6 (Conv2D)	(None, 32, 32, 256)	1179904
up_sampling2d_1 (UpSampling2D)	(None, 64, 64, 256)	0
conv2d_7 (Conv2D)	(None, 64, 64, 128)	295040
up_sampling2d_2 (UpSampling2D)	(None, 128, 128, 128)	0
conv2d_8 (Conv2D)	(None, 128, 128, 64)	73792
conv2d_9 (Conv2D)	(None, 128, 128, 32)	18464
conv2d_10 (Conv2D)	(None, 128, 128, 3)	867
Total params: 1,587,011		
Trainable params: 1,587,011		
Non-trainable params: 0		

Hyperparameters

- Number and makeup of hidden layers
- Cost Function
- Optimizer
- Bottleneck layer size
- Data Augmentation
- Image Size

```
Conv2D(filters=32, kernel_size = (3, 3), strides=1, padding='same', activation='selu', kernel_initializer='lecun_normal'))
```

```
deep_ae.compile(loss=tf.keras.losses.MeanSquaredError(), optimizer=tf.keras.optimizers.Adam(), metrics=[rounded_accuracy])
```

```
deep_e.add(keras.layers.Dense(1024, activation='selu', kernel_initializer='lecun_normal'))
```

```
resample_x = 128  
resample_y = 128  
resample_z = 3
```

Cost function and Model fitting

Cost function: Pixel by Pixel - Mean squared error

Result: ~ 0.015 MSE (1 is bad, 0 is good)

Optimizer: Adam optimizer (form of gradient descent)

MSE for one image:

$$\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}$$

Input



Output

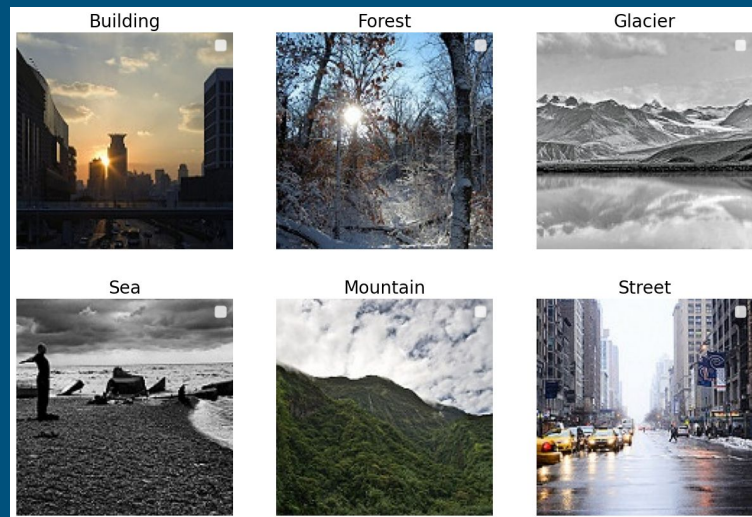
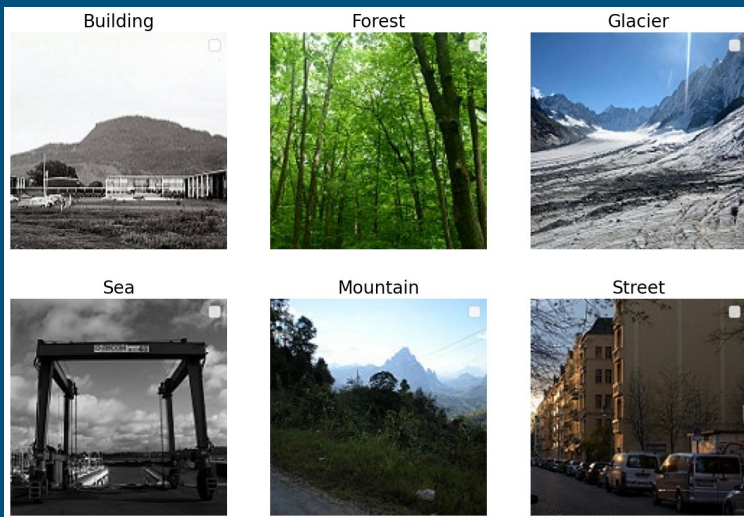


Intel Image Dataset

~14000 training instances
~3000 test instances
Balanced dataset
Image size: 150, 150, 3
6 classes of images

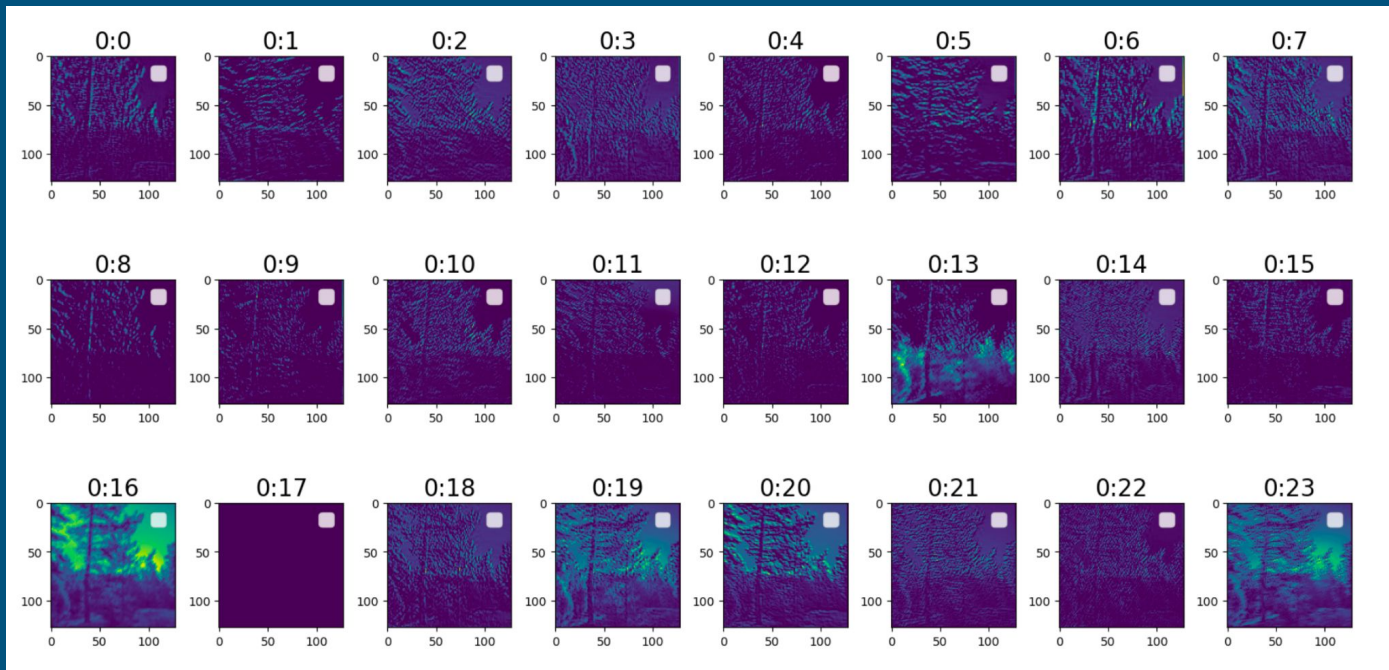

Total number of train buildings instances: 2191
Total number of train forest instances: 2271
Total number of train glacier instances: 2404
Total number of train mountain instances: 2512
Total number of train sea instances: 2274
Total number of train street instances: 2382

Total number of test buildings instances: 437
Total number of test forest instances: 474
Total number of test glacier instances: 553
Total number of test mountain instances: 525
Total number of test sea instances: 510
Total number of test street instances: 501



1st Convolution

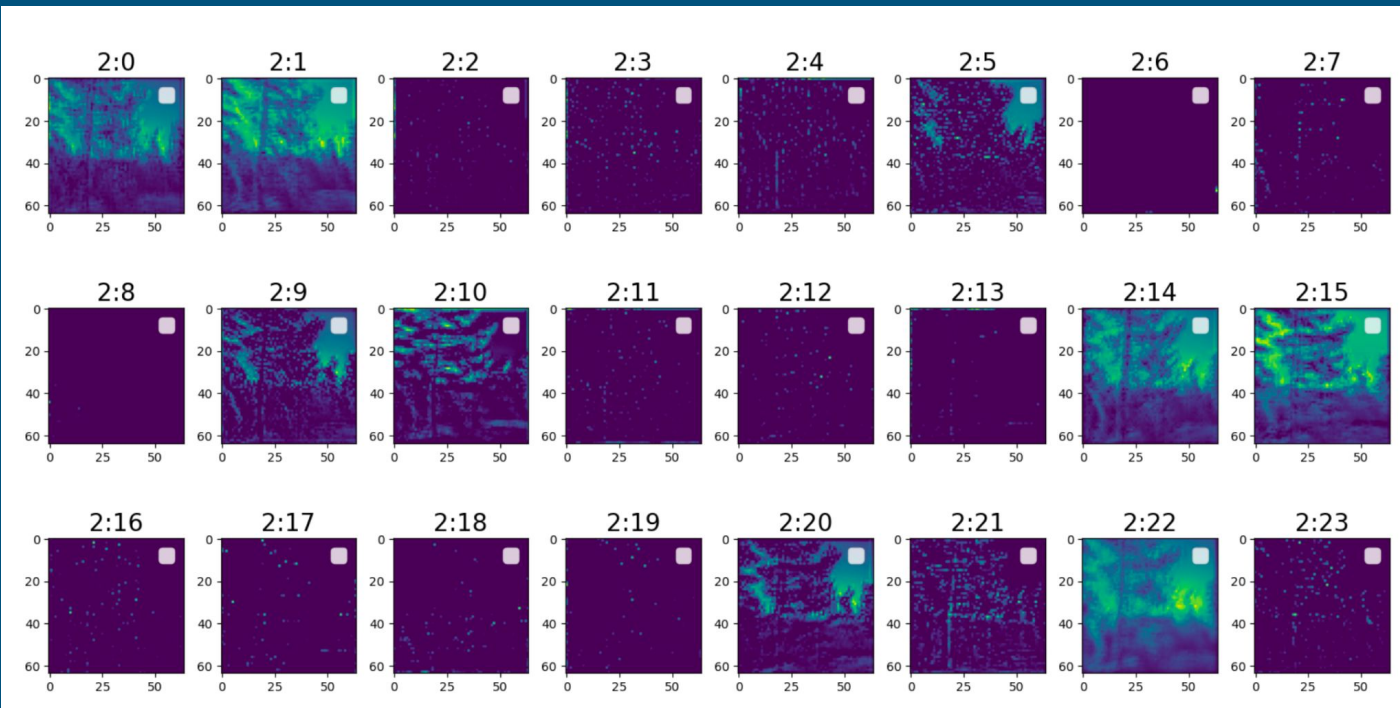

32 Filters



```
Conv2D(filters=32,kernel_size = (3, 3),strides=1,padding='same', activation='selu', kernel_initializer='lecun_normal')
```

2nd Convolution

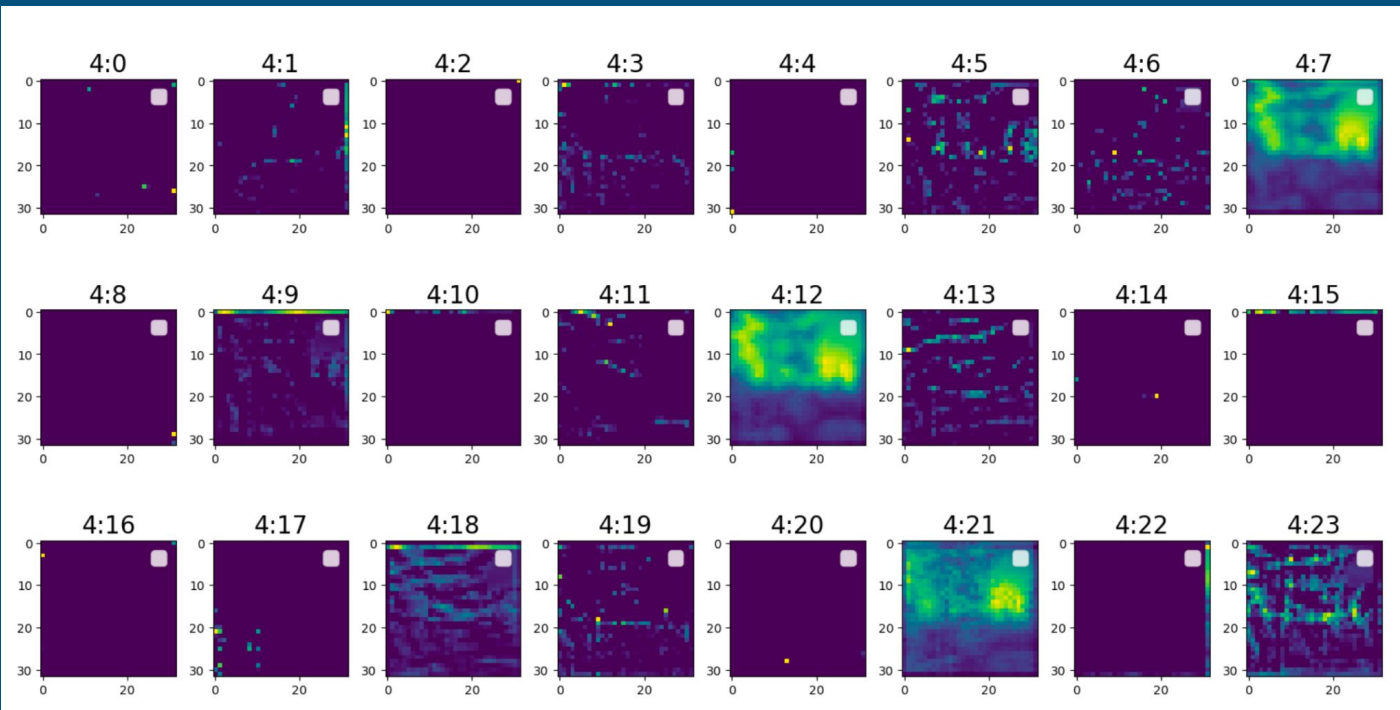

64 Filters



```
Conv2D(filters=64,kernel_size = (3, 3),strides=1,padding='same', activation='selu', kernel_initializer='lecun_normal')
```

3rd Convolution

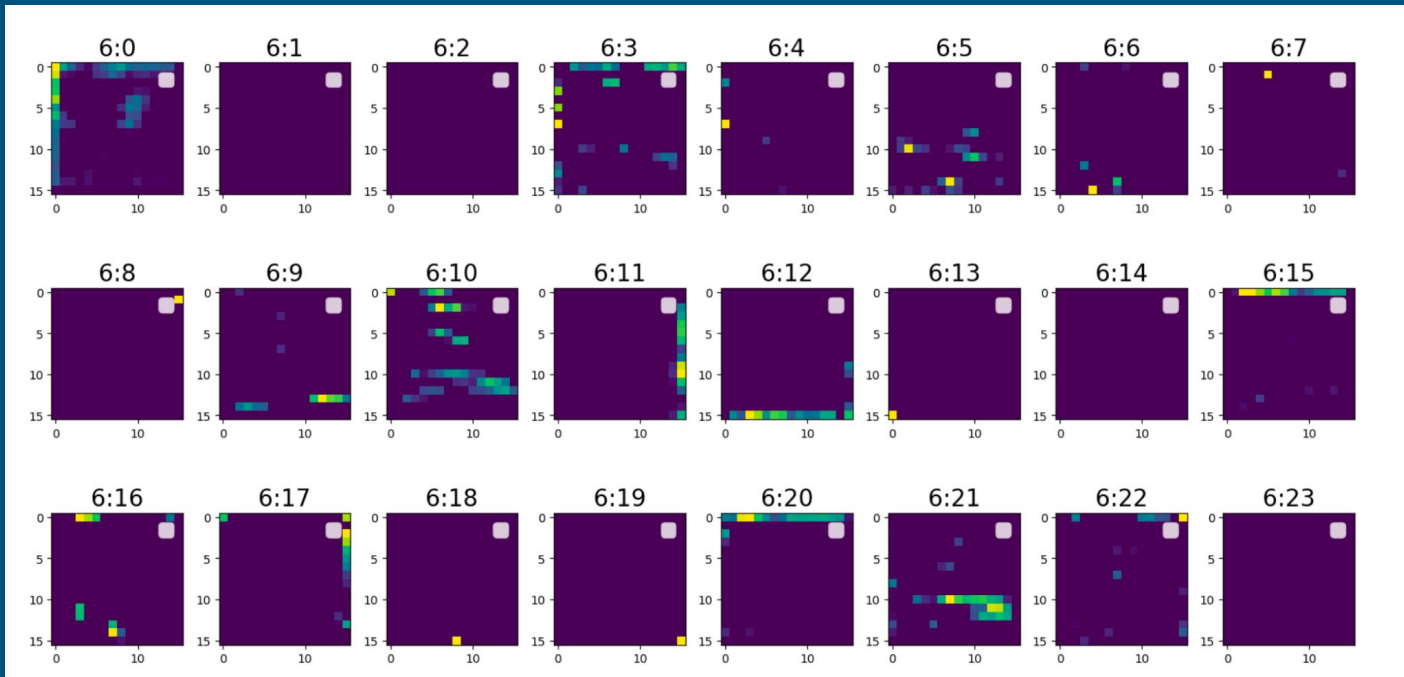

128 Filters



`Conv2D(filters=128, kernel_size = (3, 3), strides=1, padding='same', activation='selu', kernel_initializer='lecun_normal')`

4th Convolution



256 Filters



```
Conv2D(filters=256,kernel_size = (3, 3),strides=1,padding='same', activation='selu', kernel_initializer='lecun_normal')
```


Last Convolution

512 Filters



```
Conv2D(filters=512,kernel_size = (3, 3),strides=1,padding='same', activation='selu', kernel_initializer='lecun_normal')
```


Encoding



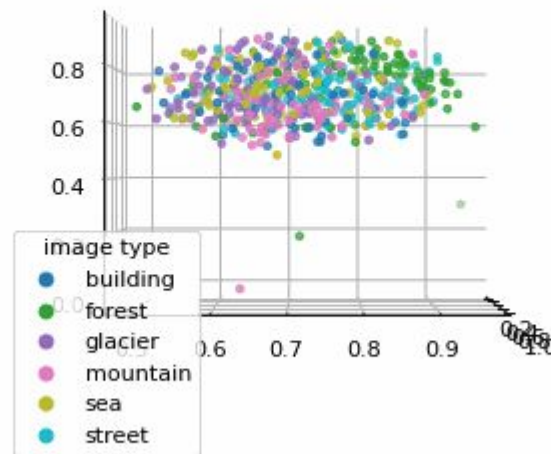
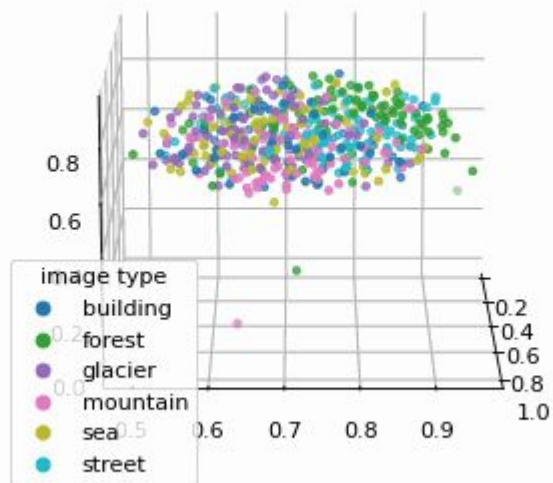
128 x 128 x 3



32 x 32 x 1

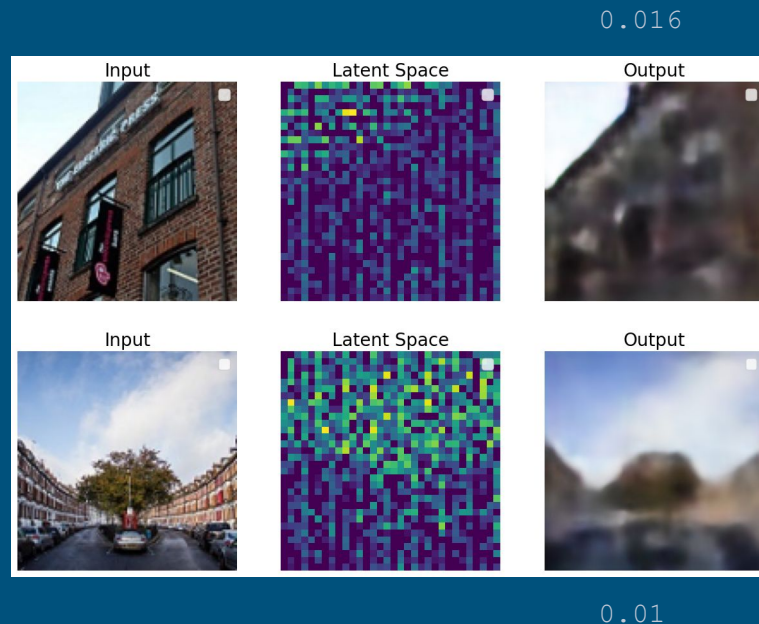
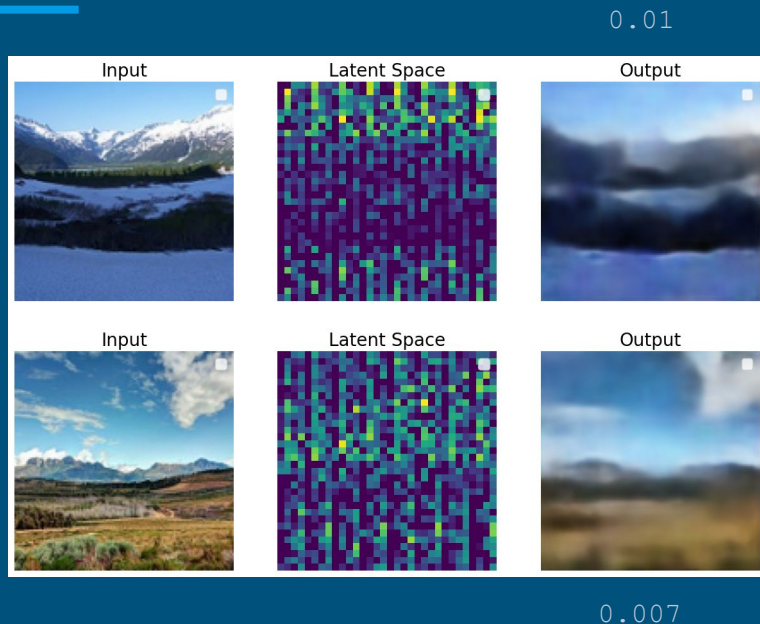


Latent Space on Validation Data



Reduced from 1024 to 3 dimensions with TSNE

Visualisation of latent space of four samples




Notice: ordering of pixels in latent space does not matter it was reshaped into this format for display purposes

Applications

- Preprocessing Data (Dimensionality Reduction)
- Image denoising
- Anomaly Detection
- Generation of new images

Image Denoising


- Example: Trained model on the MNIST handwritten digits dataset



https://miro.medium.com/max/1100/1*GFd9K-88w06YKOEqphGMng.png

Anomaly Detection

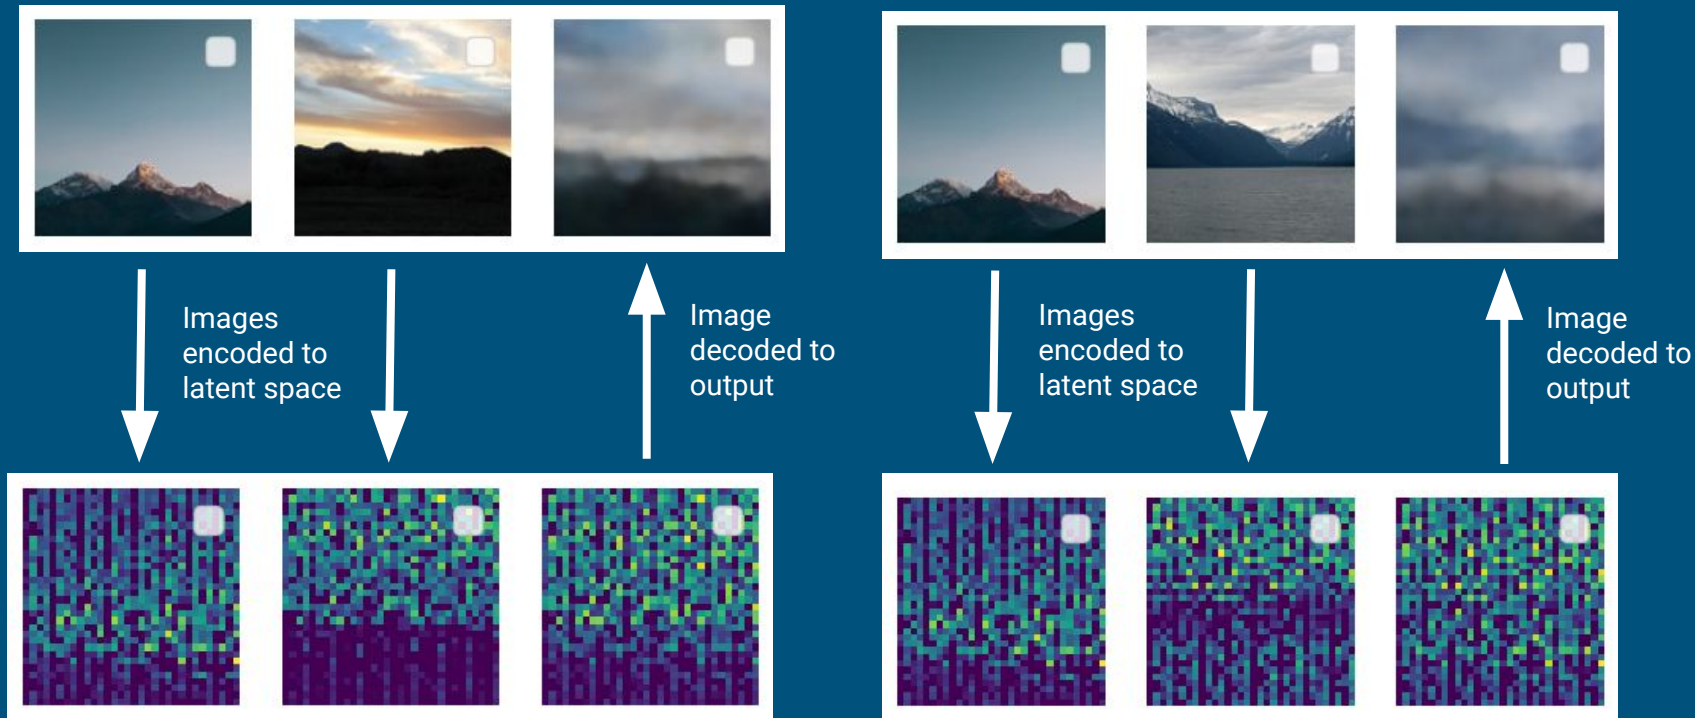
- Detect outliers in dataset
- Based on reconstruction error
- Can be passed to another model
 - (e.g. Linear Classifier, etc.)



<https://aws.amazon.com/blogs/machine-learning/deploying-variational-autoencoders-for-anomaly-detection-with-tensorflow-serving-on-amazon-sagemaker/>




Decision Boundary

Averaging two latent spaces together



Notice: ordering of pixels in latent space does not matter it was reshaped into this format for display purposes

Average of two pictures



Advantages / Disadvantages

- Advantages

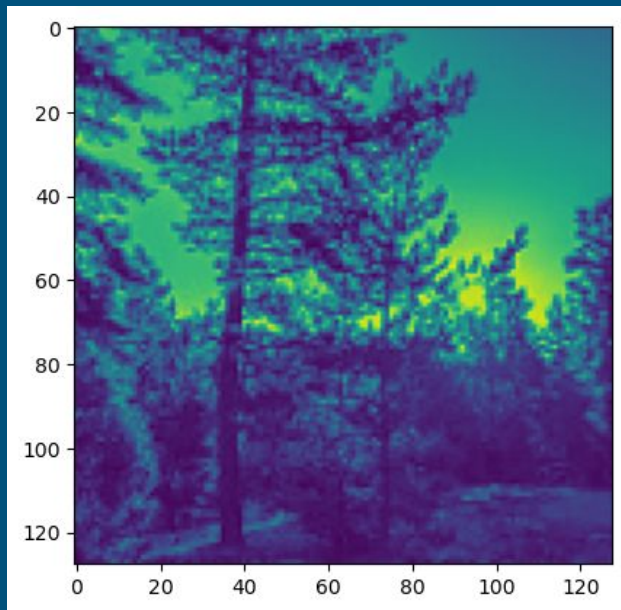
- Able to handle more complex relationships (non-linear)
- Not as 'lossy' as PCA
- Clear metric to evaluate the model
- Less computationally expensive compared to a dense network

- Disadvantages

- Slow training process (Our autoencoder took about 90 minutes for 50 epochs)
- A lot of model and hyperparameter tuning
- Encoding into very low dimensions (PCA or TSNE is better at lower dimensions <50)
- Requires lots of data


PCA compression

PCA Input Image




128 x 128 x 1

Encoding






32 x 32 x 1

PCA Inverse Transform



128 x 128 x 1

Fails while making this project



Sources

- <https://www.v7labs.com/blog/autoencoders-guide>
- <https://www.kaggle.com/datasets/puneet6060/intel-image-classification>
- <https://towardsdatascience.com/dimensionality-reduction-pca-versus-autoencoders-338fcaf3297d>
- https://keras.io/guides/transfer_learning/
- <https://aws.amazon.com/blogs/machine-learning/deploying-variational-autoencoders-for-anomaly-detection-with-tensorflow-serving-on-amazon-sagemaker/>
- <https://towardsdatascience.com/auto-encoder-what-is-it-and-what-is-it-used-for-part-1-3e5c6f017726>
- <https://www.kaggle.com/code/abdelrhmanfakhry/cnn-data-augmentation-early-stop-lr-reduction>
-
- MSOE AI Club